

DevSecOps Kubernetes Pipeline Workshop

From [@ibm](#), [@tetradeio](#), and [@controlplaneio](#)



IBM **Cloud**



Michael Hough

Developer, IBM Cloud Container Registry

Maintainer, Portieris



@molepigeon

Innovate with IBM's Cloud-Native Platform



Functions



Kubernetes



CLOUD FOUNDRY



Istio



Event
Streams



Watson



Cloudant



Databases

*Continuously delivering
innovation to your application*


Client Success Stories



MAERSK

Always-On & Automation

*IBM Cloud builds on open-source
to relieve the pains of security,
scale, software, & infrastructure
management*



Tetrade

Liam White

Software Engineer, Tetrade

Core contributor, Portieris & Istio



@liamandrewwhite





Pi Unnerup

Infrastructure Engineer, ControlPlane

OS work on Netassert, Kubesec.io



@piunnerup



Andrew Martin

Security Engineer, ControlPlane

OS work on Kubernetes & Istio



@sublimino

Preflight Checks

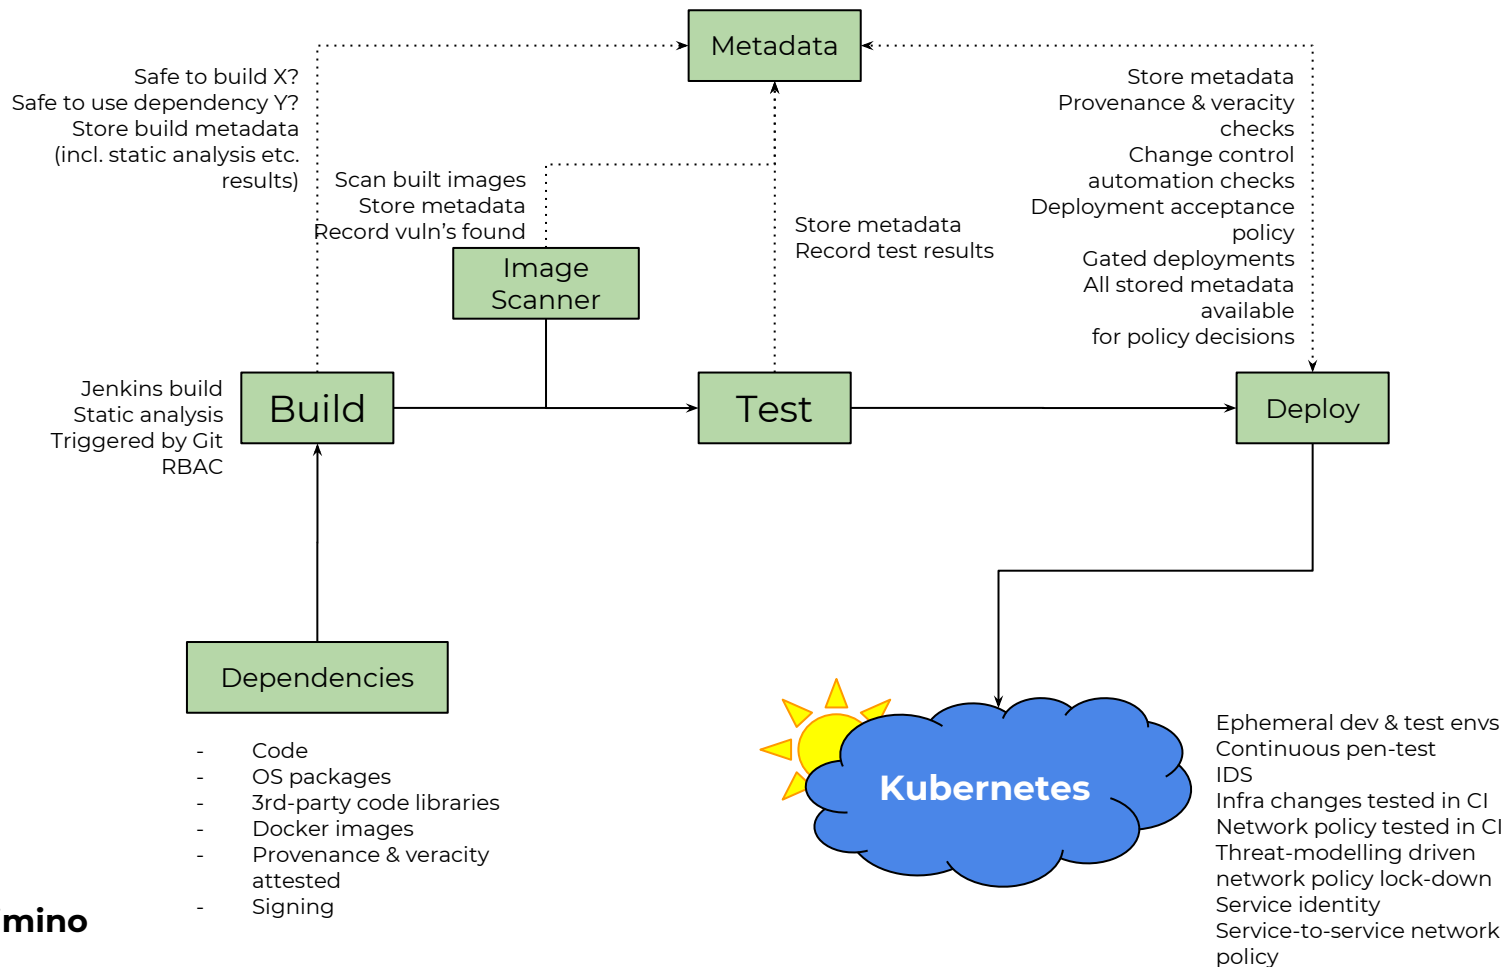
Preparation

- Navigate to <https://goo.gl/DKXRnb> and follow the instructions
 - Ensure minikube is running (tested on v1.10+)
 - `docker pull sublimino/alpine-base:insecure`
 - Ensure you can `kubectl get pods`
 - Clone the repo locally
 - Run through **00-Prerequisites** and **01-Installing-Harbor**

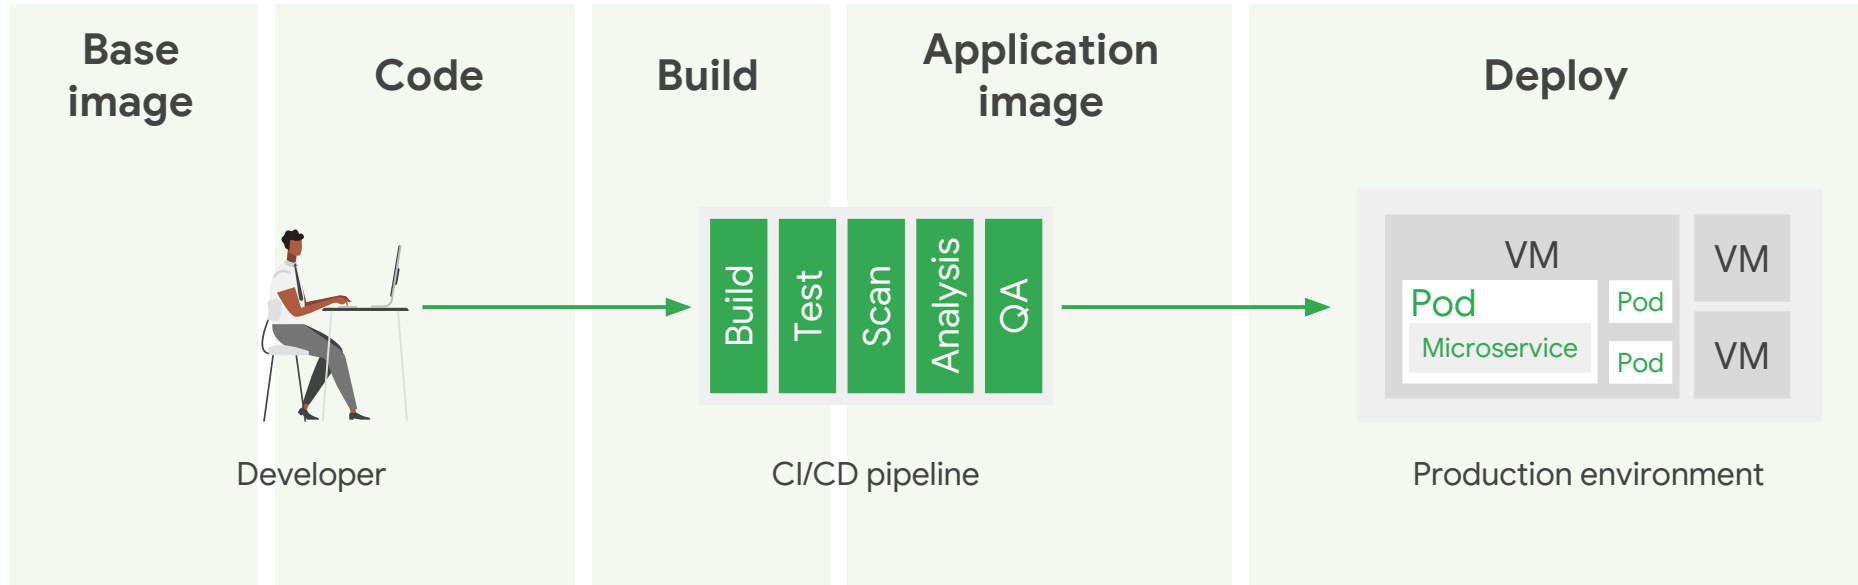
Secure Pipelines




Secure Cloud-Native Delivery




Stages of the CDLC (Container Delivery Lifecycle)



A full-body image of Darth Vader from Star Wars, standing with his arms outstretched and pointing his right index finger towards the viewer. He is wearing his iconic black helmet and suit. The background is a blurred, greyish-blue.

“I find your lack of security
disturbing.”



Open-source supply chain today

Base image

Images: Docker
Distribution (Hub)



Code

Updates: TUF,
Notary



Build

**Pipeline
metadata:**
Grafeas, in-toto



Application image

**Vulnerability
scanning:** Clair,
Micro Scanner,
Anchore Open
Source Engine



Deploy


**Admission
control:** K8s
admission
controllers, Kritis,
Portieris



Open-source supply chain today

Base image

Images: Docker
Distribution (Hub)



Code

Updates: TUF,
Notary



Build

**Pipeline
metadata:**
Grafeas, in-toto



Application image

**Vulnerability
scanning:** Clair,
Micro Scanner,
Anchore Open
Source Engine



Deploy

**Admission
control:** K8s
admission
controllers, Kritis,
Portieris



Build Flow

- Build image (base image from Docker Hub)
- Assert absence of vulnerabilities in image (Harbor)
- Cryptographically sign image for later verification
- Push image to container registry
- Attempt to deploy image to cluster
- Verify image has been signed with an admission controller
- Reject images that have not followed due process and organisational policy

01 - Installing Harbor

Harbor

- Container image registry (a “self-hosted Docker hub”)
- New CNCF project
- Capable of running inside a cluster for inception-esque self-referential image pulls




Harbor



Notary

Cryptographic
image signing




**Docker
Distribution**
Container registry



clair

Clair
Image vulnerability
scanning

Vulnerable Images





What Can Image Scanning Detect?



- This depends upon the depth of the tool
- Some will just scan installed operating system package manager versions
- Others will check filesystem permissions for all entities, extra binaries, secrets, policies etc.


Image vulnerability scanning approaches



<https://sysdig.com/blog/container-security-docker-image-scanning/>

- Components to scan:
package-level vs. code-level
 - OS packages
 - App library packages
 - JARs, WARs, TARs, etc.
 - Malware
 - Misconfigurations, e.g., secrets
- Scan type
 - Layer-by-layer
 - UnionFS top layer only

Clair vs. MicroScanner vs. Anchore



Scanning depth

OS covered

Maintainer



Packages



Packages

anchore

Packages, files,
software artifacts

Alpine, CentOS,
Debian, Oracle
Linux, RHEL, Ubuntu

CoreOS

Aqua Security

Anchore

02 - Vulnerability Scanning

Part 2 - recap

- Firstly - we couldn't deploy anything! Harbor will not allow us to pull vulnerable images
- This was the intended consequence of attempting to ship a CVE-laden image to production
- CVEs are a likely way for an attacker to bring their assault on your systems
- Never ship CVEs to production

Notary





Daemon

Digest for ubuntu:latest,
please!



Content for ubuntu@12345,
please!



Registry

12345

<stuff>





Daemon

Digest for ubuntu:latest,
please!

I trust Bob...

And that's his digital
signature!

Content for ubuntu@12345,
please!



Notary

12345, and it's signed by
Alice, Bob and Charlie




Registry

<stuff>



Admission control

Extensible Admission Controllers



<http://blog.kubernetes.io/2018/01/extensible-admission-is-beta.html>



P**RTIERIS**




image: **ibmcom/portieris:0.5.1**

image: **ibmcom/portieris@sha256:19b6e9df327....**


03 - Image Signing

Part 3 - was that enough?

- Scanning for vulnerabilities is important - but only makes any sense if that same image is deployed to production
- Asserting that the image that runs in production contains what you think it does is another basic security precaution that is too-often overlooked
- This security measure can prevent the compromise of access to your container registry from compromising production

Kubesecc

kubesecc.io - risk score for K8S YAML



index

```
.metadata .annotations
'container.seccomp.security.alpha.kubernetes.io/pod'
.metadata .annotations
'seccomp.security.alpha.kubernetes.io/pod'
.spec.template.spec.hostIPC
.spec.template.spec.hostNetwork
.spec.template.spec.hostPID
Service Accounts
containers[].resources.limits.cpu
containers[].resources.limits.memory
containers[].resources.requests.cpu
containers[].resources.requests.memory
containers[].securityContext.capabilities.add | index("SYS_ADMIN")
containers[].securityContext.capabilities.drop | index("ALL")
containers[].securityContext.privileged == true
containers[].securityContext.readOnlyRootFilesystem == true
containers[].securityContext.runAsNonRoot == true
```

kubesecc.io > index

[Edit this page](#)

KUBESEC.IO

- .metadata .annotations "container.seccomp.security.alpha.kubernetes.io/pod"
- .metadata .annotations "seccomp.security.alpha.kubernetes.io/pod"
- .spec.template.spec.hostIPC
- .spec.template.spec.hostNetwork
- .spec.template.spec.hostPID
- Service Accounts
- containers[].resources.limits.cpu
- containers[].resources.limits.memory
- containers[].resources.requests.cpu
- containers[].resources.requests.memory
- containers[].securityContext.capabilities.add | index("SYS_ADMIN")
- containers[].securityContext.capabilities.drop | index("ALL")
- containers[].securityContext.privileged == true
- containers[].securityContext.readOnlyRootFilesystem == true
- containers[].securityContext.runAsNonRoot == true
- containers[].securityContext.runAsUser > 10000
- securityContext.capabilities
- .metadata .annotations "container.apparmor.security.beta.kubernetes.io/nginx"
- .spec.volumeClaimTemplates[].spec.accessModes | index("ReadWriteOnce")
- .spec.volumeClaimTemplates[].spec.resources.requests.storage

FUTHER READING

- <http://blog.kubernetes.io/2016/08/security-best-practices-kubernetes-deployment.html>



kubesecc.io - example insecure pod

```
{
  "score": -30,
  "scoring": {
    "critical": [{
      "selector": "containers[] .securityContext .privileged == true",
      "reason": "Privileged containers can allow almost completely unrestricted host access"
    }],
    "advise": [{
      "selector": "containers[] .securityContext .runAsNonRoot == true",
      "reason": "Force the running image to run as a non-root user to ensure least privilege"
    }, {
      "selector": "containers[] .securityContext .capabilities .drop",
      "reason": "Reducing kernel capabilities available to a container limits its attack surface",
      "href": "https://kubernetes.io/docs/tasks/configure-pod-container/security-context/"
    }],
    ...
  }
}
```

04 - More Admission Control

Part 4 - minimum viable security

- We have
 - Verified the contents of an image are not insecure
 - Signed the image to confirm we have tested it
 - Prevented unsigned images from being deployed to production
- These are the building blocks of a secure pipeline
 - But only focus on the contents of the image and not its runtime configuration
- PodSecurityPolicy and NetworkPolicy should be use to limit the behaviour of the application at runtime
- Further admission controllers can be added to enhance security

Threat Model

- Attacks wholly or partially mitigated:
 - Container image and application supply chain with known CVEs
 - Theft of users' container registry credentials
 - Some build server compromises
- Extant risk:
 - Compromised user or insider threat
 - Zero day vulnerabilities
 - ...the rest of the Kubernetes attack surface!

Summary